

## A fast algorithm for matrix balancing

PHILIP A. KNIGHT\*

*Department of Mathematics and Statistics, University of Strathclyde, 26 Richmond Street,  
Glasgow G1 1XH, Scotland, UK*

\*Corresponding author: p.a.knight@strath.ac.uk

AND

DANIEL RUIZ

*INPT-ENSEEIH, 2, rue Charles Camichel, Toulouse, France*

Email: daniel.ruiz@enseeiht.fr

[Received on 5 July 2011; revised on 7 June 2012; accepted on 7 June 2012]

As long as a square non-negative matrix  $A$  has total support then it can be balanced, that is, we can find a diagonal scaling of  $A$  that has row and column sums equal to one. A number of algorithms have been proposed to achieve the balancing, the most well known of these being Sinkhorn–Knopp. In this paper, we derive new algorithms based on inner–outer iteration schemes. We show that Sinkhorn–Knopp belongs to this family, but other members can converge much more quickly. In particular, we show that while stationary iterative methods offer little or no improvement in many cases, a scheme using a preconditioned conjugate gradient method as the inner iteration converges at much lower cost (in terms of matrix–vector products) for a broad range of matrices; and succeeds in cases where the Sinkhorn–Knopp algorithm fails.

**Keywords:** matrix balancing; Sinkhorn–Knopp algorithm; doubly stochastic matrix; conjugate gradient iteration.

### 1. Introduction

For at least 70 years, scientists in a wide variety of disciplines have attempted to balance square non-negative matrices by applying diagonal scalings. That is, given  $A \in \mathbb{R}^{n \times n}$ ,  $A \geq 0$ , find diagonal matrices  $D_1$  and  $D_2$  so that all the rows and columns of  $P = D_1 A D_2$  sum to one. In statistical applications (and informally elsewhere)  $P$  is called doubly stochastic. Motivations for achieving this balance include interpreting economic data (Bacharach, 1970), understanding traffic circulation (Lamond & Stewart, 1981), assigning seats fairly after elections (Balinski, 2008), matching protein samples (Daszykowski *et al.*, 2009) and ordering nodes in a graph (Knight, 2008). There are also similar applications when  $A$  is not necessarily non-negative, such as preconditioning sparse matrices (Livne & Golub, 2004), where we call a matrix balanced if the one norms of each row and column is one. In all of these applications, one of the main methods considered is SK (the algorithm has been given many different names, Sinkhorn–Knopp is the name usually adopted by the linear algebra community; we will refer to it as SK in the rest of this paper). This is an iterative process that attempts to find  $D_1$  and  $D_2$  by alternately normalizing columns and rows in a sequence of matrices starting with  $A$ . Convergence conditions for this algorithm are well known: if  $A$  has total support (defined in Section 3) then the algorithm will converge linearly

with asymptotic rate equal to the square of the subdominant singular value of  $P$  (Sinkhorn & Knopp, 1967; Soules, 1991; Knight, 2008).

Clearly, in some cases the convergence will be painfully slow. The principal aim of this paper is to derive some new algorithms for the matrix-balancing problem with an eye on speed, especially for large systems. First, we look at a simple Newton method for symmetric matrices, closely related to a method proposed in Khachiyan & Kalantari (1992), for positive definite (but not necessarily non-negative) matrices. We will show that as long as Newton's method produces a sequence of positive iterates, the Jacobians we generate will be positive semidefinite and that this is also true when we adapt the method to cope with nonsymmetric matrices.

To apply Newton's method exactly we need to solve a linear system at each step, and this is usually prohibitively expensive. We therefore look at iterative techniques for approximating the solution at each step. First, we look at splitting methods and we see that SK is a member of this family of methods, as is the algorithm proposed in Livne & Golub (2004). We give an asymptotic bound on the (linear) rate of convergence of these methods. For symmetric matrices, we can get significant improvement over SK. Unfortunately, we show this is not true for nonsymmetric matrices.

Next we look at a preconditioned conjugate gradient method for solving the linear system. We discuss implementation details and show that asymptotically, quadratic convergence is achievable, even for nonsymmetric matrices. By implementing a box constraint we ensure that our iterates retain positivity and we demonstrate the reliability and speed of our method in tests. The algorithm, the core of which takes up less than 40 lines of MATLAB code, is given in the Appendix.

A number of authors have presented alternative techniques for balancing matrices that can converge faster than SK. For example, Parlett & Landis (1982) look at some simple ways of trying to accelerate the convergence by focusing on reducing statistics such as the standard deviation between row sums of the iterates. In certain cases, they show that great improvement is possible, suggesting that the rate of convergence is not dependent on the singular values of  $P$ ; but they also give examples where their alternatives perform significantly worse. We include a comparison of one of these algorithms against our proposed approach in Section 6. Linial *et al.* (2000) use a similar approach to Parlett & Landis (1982) in the context of estimating matrix permanents, although their upper bound on iteration counts ( $\mathcal{O}(n^7)$ ) makes them distinctly unappealing for large problems.

If  $A$  is non-negative then our new algorithm can be used without directly accessing the matrix: we simply need to be able to form matrix–vector products. If we want to balance the norms of general matrices we need to be able to form products involving  $|A|x$ . In the context of preconditioning, we may need only an approximate scaling of  $|A|$ . Matrix-free algorithms have been developed to find approximate scalings (Bradley, 2010).

A completely different approach is to view matrix balancing as an optimization problem. There are many alternative formulations, perhaps the first being that of Marshall & Olkin (1968) who show that balancing is equivalent to minimizing the bilinear form  $x^T A y$  subject to the constraints  $\Pi x_i = \Pi y_i = 1$ . However, the experimental results we have seen for optimization techniques for balancing (Marshall & Olkin, 1968; Schneider, 1990; Balakrishnan *et al.*, 2004) suggest that these methods are not particularly cheap to implement.

## 2. Newton's method

Let  $\mathcal{D} : \mathbb{R}^n \rightarrow \mathbb{R}^{n \times n}$  represent the operator that turns a vector into a diagonal matrix,  $\mathcal{D}(x) = \text{diag}(x)$ , and let  $e$  represent a vector of ones. Then to balance a non-negative matrix,  $A$ , we need to find positive

vectors  $r$  and  $c$  such that

$$\mathcal{D}(r)A\mathcal{D}(c)e = \mathcal{D}(r)Ac = e, \quad \mathcal{D}(c)A^T r = e. \quad (2.1)$$

Rearranging these identities gives

$$c = \mathcal{D}(A^T r)^{-1} e, \quad r = \mathcal{D}(Ac)^{-1} e,$$

where  $\mathcal{D}(x)^{-1}$  is the inverse of  $\text{diag}(x)$ , and one way of writing SK (Kalantari & Khachiyan, 1996; Knight, 2008) is as

$$c_{k+1} = \mathcal{D}(A^T r_k)^{-1} e, \quad r_{k+1} = \mathcal{D}(Ac_{k+1})^{-1} e. \quad (2.2)$$

If  $A$  is symmetric then (2.1) can be simplified. To achieve balance we need a vector  $x_*$  that satisfies

$$f(x_*) = \mathcal{D}(x_*)Ax_* - e = 0. \quad (2.3)$$

This leads to the iterative step

$$x_{k+1} = \mathcal{D}(Ax_k)^{-1} e. \quad (2.4)$$

Exactly the same calculations are performed as for (2.2), with  $r_k$  corresponding to  $x_{2k-1}$  and  $c_{k+1}$  to  $x_{2k}$  (Knight, 2008).

An obvious alternative to SK is to solve (2.3) with Newton's method. Differentiating  $f(x)$  gives

$$J(x) = \frac{\partial}{\partial x} (\mathcal{D}(x)Ax - e) = \mathcal{D}(x)A + \mathcal{D}(Ax),$$

a result that is easily confirmed by componentwise calculation or with tensor algebra; hence Newton's method can be written as

$$x_{k+1} = x_k - (\mathcal{D}(x_k)A + \mathcal{D}(Ax_k))^{-1} (\mathcal{D}(x_k)Ax_k - e).$$

We can rearrange this equation to get

$$(\mathcal{D}(x_k)A + \mathcal{D}(Ax_k))x_{k+1} = \mathcal{D}(Ax_k)x_k + e,$$

so,

$$\begin{aligned} (A + \mathcal{D}(x_k)^{-1}\mathcal{D}(Ax_k))x_{k+1} &= \mathcal{D}(x_k)^{-1}(\mathcal{D}(Ax_k)x_k + e) \\ &= Ax_k + \mathcal{D}(x_k)^{-1}e, \end{aligned} \quad (2.5)$$

and we can set up each Newton iteration by performing some simple vector operations and updating the diagonal on the left-hand side. This matrix plays a key role in our analysis and we introduce the notation

$$\mathcal{A}_k = A + \mathcal{D}(x_k)^{-1}\mathcal{D}(Ax_k).$$

Note that this matrix inherits the symmetry of  $A$ . An algorithm to balance  $A$  can be implemented by applying one's linear solver of choice iteratively to (2.5). In practical applications, it makes sense to apply an inner-outer iteration scheme. In Sections 4 and 5 we look at some efficient ways of doing this.

The idea of using Newton's method to solve the scaling problem is not new and was first proposed in [Khachiyan & Kalantari \(1992\)](#). Instead of (2.3), the equivalent equation

$$Ax - \mathcal{D}(x)^{-1}e = 0 \quad (2.6)$$

is considered. The authors were interested in the problem of scaling symmetric positive definite (SPD) matrices, rather than non-negative matrices. In this case, the scaled matrix need not be doubly stochastic. While its row and column sums are one, it may contain negative entries. The authors did not consider the practical application of their algorithm, although the methods we look at in this paper can be adapted to work for their formulation. In Section 5 we explain why our approach leads to faster convergence.

Newton's method is also used as a method for solving the balancing problem for symmetric matrices in [Livne & Golub \(2004\)](#). Here the authors work with

$$\left(I - \frac{ee^T}{n}\right) \mathcal{D}(Ax)x = 0 \quad (2.7)$$

instead of (2.3). They then use a Gauss–Seidel Newton method to solve the problem and show that this approach can give significant improvements over SK. We will develop this idea in Section 4.

Yet another formulation of (2.4) can be found in [Fürer \(2004\)](#), where the author suggests the resulting equation is solved by Newton's method. However no attempt is made to implement the suggested algorithm and the fact that only the right-hand side changes in the linear system that is solved at each step suggests that rapid convergence is unlikely in general.

Our new algorithm is not restricted to symmetric matrices. Note that if  $A$  is nonsymmetric, one can use (2.4) on

$$S = \begin{bmatrix} 0 & A \\ A^T & 0 \end{bmatrix},$$

each iteration with  $S$  representing two SK steps. We can replace  $A$  in (2.5) with  $S$ . The resulting linear systems are singular but, as we will see later, this does not prevent convergence.

### 3. Properties of $\mathcal{A}_k$

In order that we can solve the balancing problem efficiently, in particular when  $A$  is large and sparse, we will use iterative methods to approximately solve the linear system in (2.5). There are a number of possibilities to choose because  $\mathcal{A}_k$ , the matrix on the left-hand side of the expression, is symmetric positive semidefinite. This is a consequence of the following result (which does not require symmetry).

**THEOREM 3.1** Suppose that  $A \in \mathbb{R}^{n \times n}$  is non-negative and  $y \in \mathbb{R}^n$  is positive. Let  $D = \mathcal{D}(Ay)\mathcal{D}(y)^{-1}$ . Then for all eigenvalues  $\lambda$  of  $A + D$ ,  $\operatorname{Re}(\lambda) \geq 0$ .

*Proof.* Note that  $A + D$  is similar to

$$\begin{aligned} \mathcal{D}(y)^{-1}(A + D)\mathcal{D}(y) &= \mathcal{D}(y)^{-1}A\mathcal{D}(y) + \mathcal{D}(y)^{-1}\mathcal{D}(Ay) \\ &= \mathcal{D}(y)^{-1}A\mathcal{D}(y) + \mathcal{D}(\mathcal{D}(y)^{-1}A\mathcal{D}(y))e \\ &= B + \mathcal{D}(Be), \end{aligned}$$

where  $B = \mathcal{D}(y)^{-1}A\mathcal{D}(y)$ . Now  $Be$  is simply the vector of row sums of  $B$  and so adding this to the diagonal of  $B$  gives us a diagonally dominant matrix. Since the diagonal is non-negative, the result follows from Gershgorin's theorem.  $\square$

If  $A$  has no zero rows and a nonzero entry on its diagonal then at least one of the rows of  $B + \mathcal{D}(Be)$  will be strongly diagonally dominant and, by a theorem of Taussky (1948), it will be nonsingular. We can also ensure that  $\mathcal{A}_k$  is nonsingular by imposing conditions on the connectivity between the nonzeros in  $A$ . We can establish the following result.

**THEOREM 3.2** Suppose that  $A \in \mathbb{R}^{n \times n}$  is non-negative. If  $A$  is fully indecomposable then  $\mathcal{A}_k$  is nonsingular.

Recall that a matrix  $A$  is fully indecomposable if it is impossible to find permutation matrices  $P$  and  $Q$  such that

$$PAQ = \begin{bmatrix} A_1 & 0 \\ A_2 & A_3 \end{bmatrix}$$

with  $A_1$  and  $A_2$  square (a generalization of irreducibility).

The proof of Theorem 3.2 is postponed to the end of the section. However in many cases we will not be able to satisfy its conditions. For example, if  $A$  is nonsymmetric and we use Newton's method to balance

$$S = \begin{bmatrix} 0 & A \\ A^T & 0 \end{bmatrix},$$

then  $S$  is not fully indecomposable. In fact it is straightforward to show that in this case the linear systems we have to solve are singular.

We would also like to use our new algorithms on any non-negative matrix for which balancing is possible, namely, any matrix which has total support.  $A$  has **total support** if it is nonzero and all its nonzero elements lie on a positive diagonal. That is, if  $a_{ij} \neq 0$  we can find a permutation,  $B$ , of the rows of  $A$  which puts  $a_{ij}$  on the leading diagonal and for which  $|b_{kk}| > 0$  for all  $k$ . A matrix has **support** if it has at least one positive diagonal. Matrices that have total support but that are not fully indecomposable also lead to singular systems in (2.5). Singularity in these cases is not problematic as the systems are consistent. As Theorem 3.5 shows, we can go further, and we can solve the systems whenever  $A$  has support.

We need some preliminary results.

**LEMMA 3.3** Suppose that  $A \in \mathbb{R}^{n \times n}$  is a symmetric non-negative matrix with support. Then there is a permutation  $P$  such that  $PAP^T$  is a block diagonal matrix in which all the diagonal blocks are irreducible.

*Proof.* We show this by induction on  $n$ . Clearly it is true if  $n = 1$ . Suppose  $n > 1$  and our hypothesis is true for all matrices of dimension smaller than  $n$ . If  $A$  is irreducible there is nothing to prove, otherwise we can find a permutation  $Q$  such that

$$QAQ^T = \begin{bmatrix} A_1 & C \\ 0 & A_2 \end{bmatrix} = \begin{bmatrix} A_1 & 0 \\ 0 & A_2 \end{bmatrix},$$

where  $C = 0$  by symmetry. Since  $A$  has support so does  $QAQ^T$  and hence  $A_1$  and  $A_2$  must each have support too, and we can apply our inductive hypothesis.  $\square$

LEMMA 3.4 Suppose that  $A \geq 0$  has support and let  $B = A + \mathcal{D}(Ae)$ . If  $A$  is either symmetric or irreducible then the null space of  $B$  is orthogonal to  $e$  and the null space has a basis whose elements can be permuted into the form

$$\begin{bmatrix} e \\ -e \\ 0 \end{bmatrix}.$$

(In the case that the null space has dimension one, the zero component can be omitted.)

*Proof.* If  $B$  is nonsingular we have nothing to prove, so we assume it is singular (and hence by Taussky's theorem (Taussky, 1948),  $A$  has an empty main diagonal).

From Kolotilina (2003, Theorem 2.3 and Remark 2.9) we can infer that if  $B$  is irreducible then the null space is one-dimensional and, since it is weakly diagonally dominant, all components are of equal modulus. All that remains is to show that there is an equal number of positive and negative components in elements from the null space and hence they can be permuted into the required form.

We choose a permutation  $i_1, i_2, \dots, i_n$  of  $1, \dots, n$  such that for  $1 \leq j \leq n$ ,  $b_{ij_{j+1}} \neq 0$ , where  $i_{n+1} = i_1$ . Such a permutation exists because  $A$  has support and has an empty main diagonal. Now suppose that  $x$  is in the null space of  $B$ . Since  $B$  is diagonally dominant, the sign of  $x_i$  must be opposite to that of  $x_j$  whenever  $b_{ij} \neq 0$  ( $i \neq j$ ). By construction,  $x_{i_j} = -x_{i_{j+1}}$  for  $1 \leq j \leq n$ . This is only possible if  $n$  is even, in which case  $x^T e = 0$ , as required.

If  $B$  is not irreducible but is symmetric then (by the previous lemma)  $PBP^T = \text{diag}(B_1, B_2, \dots, B_k)$  where the  $B_i$  are irreducible. Since the null space of  $B$  is formed from a direct sum of the null space of the diagonal blocks, it too must be orthogonal to  $e$  and a basis of the requisite form clearly exists.  $\square$

THEOREM 3.5 Suppose that  $A \in \mathbb{R}^{n \times n}$  is a symmetric non-negative matrix with support and that  $y > 0$ . The system

$$(A + \mathcal{D}(Ay)\mathcal{D}(y)^{-1})z = Ay + \mathcal{D}(y)^{-1}e$$

is consistent.

*Proof.* Let  $A_y = \mathcal{D}(y)A\mathcal{D}(y)$ . Since  $(A_y + \mathcal{D}(A_ye))$  is symmetric, elements orthogonal to its null space must lie in its range, so as a consequence of the previous lemma, we can find a vector  $d$  such that

$$(A_y + \mathcal{D}(A_ye))d = e.$$

Let  $z = y/2 + \mathcal{D}(y)d$ . Then

$$\begin{aligned} (A + \mathcal{D}(y)^{-1}\mathcal{D}(Ay))z &= \frac{1}{2}(A + \mathcal{D}(Ay)\mathcal{D}(y)^{-1})y + \mathcal{D}(y)^{-1}(\mathcal{D}(y)A\mathcal{D}(y)d + \mathcal{D}(Ay)\mathcal{D}(y)d) \\ &= Ay + \mathcal{D}(y)^{-1}(A_y + \mathcal{D}(A_ye))d \\ &= Ay + \mathcal{D}(y)^{-1}e. \end{aligned}$$

$\square$

While we have proved that Newton's step will converge if  $x_k > 0$ , we have not shown that  $x_{k+1} > 0$ . In fact, it need not be, and this will be an important consideration for us in developing balancing methods in the later sections.

To finish this section, we prove Theorem 3.2.

*Proof of Theorem 3.2.* Let  $B = \mathcal{D}(x_k)^{-1}A\mathcal{D}(x_k)$  and consider  $B + \mathcal{D}(Be)$ . Suppose that this matrix is singular and  $v$  lies in its null space. We can apply Lemma 3.4, and we know that exactly half of its

components are positive and half are negative. Suppose that  $P$  permutes  $v$  so that the first half of the entries of  $Pv$  are positive. Then

$$P(B + \mathcal{D}(Be))P^T = \begin{bmatrix} D_1 & B_1 \\ B_2 & D_2 \end{bmatrix},$$

where  $D_1$  and  $D_2$  are diagonal: otherwise, diagonal dominance would force some of the entries of  $(B + \mathcal{D}(Be))v$  to be nonzero. Hence

$$PBP^T = \begin{bmatrix} 0 & B_1 \\ B_2 & 0 \end{bmatrix}.$$

But such a matrix is not fully indecomposable, contradicting our hypothesis.  $\square$

#### 4. Stationary iterative methods

If we only want to solve (2.5) approximately, the simplest class of methods to consider is that of stationary iterative methods, in particular splitting methods. As  $\mathcal{A}_k$  is symmetric positive semidefinite we know that many of the common splitting methods will converge.

This approach is used in [Livne & Golub \(2004\)](#) to solve the formulation of the balancing problem given in (2.7). Here, the authors use Gauss–Seidel to solve the Newton step and demonstrate that the rate of convergence is frequently faster than that of SK. Suppose that the symmetric matrix  $A$  can be balanced so that  $P = DAD$  has equal row and column sums. Following the standard analysis for splitting methods, they give a bound on the rate of convergence in the neighbourhood of the solution, relating it to the modulus of the subdominant eigenvalue of  $L^{-1}U$ , where  $L$  is the lower-triangular part and  $U$  the strictly upper-triangular part of  $P + I - 2ee^T/n$ . Note that the bound is not in terms of  $\rho(L^{-1}U)$  as the matrix on the left-hand side of (2.7) is singular.

To measure the rate of convergence for an iterative process  $\mathcal{I}$  we use the root-convergence factor ([Ortega & Rheinboldt, 1970](#), § 9.2.1)

$$R(x) = \sup \left\{ \limsup_{k \rightarrow \infty} \|x_k - x\|^{1/k} \mid \{x_k\} \in C \right\},$$

where  $C$  is the set of all sequences generated by  $\mathcal{I}$  which converge to  $x$ .

In our main result in this section, we show that the rate of convergence  $R(x_*)$  can be ascertained for a much wider class of splitting methods. We use (2.5) rather than (2.7) as it leads to particularly simple representations of the methods.

Suppose that we use the splitting  $\mathcal{A}_k = M_k - N_k$  where  $M_k$  is nonsingular. We can then attempt to solve the balancing problem with an inner–outer iteration where the outer iteration is given by the Newton step and the inner iteration by the splitting method

$$M_k y_{j+1} = N_k y_j + Ax_k + \mathcal{D}(x_k)^{-1}e,$$

where  $y_0 = x_k$ .

Since  $\mathcal{A}_k$  and  $\mathcal{A}_{k+1}$  will differ only in their diagonals, it is usually extremely straightforward to update the splittings at any point and so we see no need to prolong the inner iteration. Indeed, if we

limit ourselves to a single inner iteration between updates, our splitting method for (2.5) becomes

$$\begin{aligned} M_k x_{k+1} &= N_k x_k + A x_k + \mathcal{D}(x_k)^{-1} e \\ &= (M_k - \mathcal{D}(x_k)^{-1} \mathcal{D}(A x_k)) x_k + \mathcal{D}(x_k)^{-1} e \\ &= M_k x_k - \mathcal{D}(A x_k) e + \mathcal{D}(x_k)^{-1} e \\ &= M_k x_k - A x_k + \mathcal{D}(x_k)^{-1} e. \end{aligned}$$

In other words, our iteration can be written as

$$x_{k+1} = x_k + M_k^{-1} (\mathcal{D}(x_k)^{-1} e - A x_k). \quad (4.1)$$

Note that if we choose  $M_k = \mathcal{D}(x_k)^{-1} \mathcal{D}(A x_k)$  in (4.1), then the resulting iteration is

$$\begin{aligned} x_{k+1} &= x_k + \mathcal{D}(x_k) \mathcal{D}(A x_k)^{-1} (\mathcal{D}(x_k)^{-1} e - A x_k) \\ &= x_k + \mathcal{D}(A x_k)^{-1} e - \mathcal{D}(x_k) e \\ &= \mathcal{D}(A x_k)^{-1} e, \end{aligned}$$

and we recover (2.4); hence SK can be neatly classified among our splitting methods.

To prove a general result we use from Ortega & Rheinboldt (1970, Theorem 10.3.1) which in our notation can be stated as follows.

**THEOREM 4.1** Let  $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$  be Gateaux differentiable in an open neighbourhood  $S_0$  of a point  $x_*$  at which  $f'$  is continuous and  $f(x_*) = 0$ . Suppose that  $f'(x) = M(x) - N(x)$  where  $M(x_*)$  is continuous and nonsingular at  $x_*$  and that  $\rho(G(x_*)) < 1$  where  $G(x) = M(x)^{-1} N(x)$ . Then for  $m \geq 1$  the iterative process  $\mathcal{I}$  defined by

$$x_{k+1} = x_k - \left( \sum_{j=1}^m G(x_k)^{j-1} \right) M(x_k)^{-1} f(x_k), \quad k = 0, 1, \dots, \quad (4.2)$$

has attraction point  $x_*$  and  $R(x_*) = \rho(G(x_*)^m)$ .

**COROLLARY 4.2** Suppose that  $A \in \mathbb{R}^{n \times n}$  is a symmetric non-negative matrix with support. Let  $x_* > 0$  be a vector such that  $P = \mathcal{D}(x_*) A \mathcal{D}(x_*)$  is stochastic. Then for the iterative process defined by (4.1),  $R(x_*) = \rho(M^{-1} N)$  where  $M - N$  is the splitting applied to  $A + \mathcal{D}(x_*)^{-2}$ .

*Proof.* The existence of  $x_*$  is guaranteed (Sinkhorn & Knopp, 1967). Let  $m = 1$ ,  $f(x) = \mathcal{D}(x) A x - e$  and let  $M(x_k) = \mathcal{D}(x_k) M_k$ . Then (4.1) can be rewritten as (4.2) and

$$G(x_k) = M(x_k)^{-1} N(x_k) = M_k^{-1} N_k.$$

So in the limit for a sequence  $\{x_k\}$  converging to  $x_*$ ,  $R(x)$  is governed by the spectral radius of the splitting of

$$\lim_{k \rightarrow \infty} \mathcal{A}_k = A + \mathcal{D}(x_*)^{-1} \mathcal{D}(A x_*) = A + \mathcal{D}(x_*)^{-2}.$$

□



For many familiar splittings,  $R(x_*)$  has an alternative representation as we can often show that  $\rho(M^{-1}N)$  is invariant under the action of diagonal scaling on  $A$ . If this is the case we can replace  $A + \mathcal{D}(x_*)^{-2}$  with

$$\mathcal{D}(x_*)(A + \mathcal{D}(x_*)^{-2})\mathcal{D}(x_*) = P + I,$$

that is, the asymptotic rate of convergence can be measured by looking at the spectral radius of the splitting matrix for  $P + I$ . Many familiar splitting methods exhibit this scaling invariance. For example, the Jacobi method, Gauss–Seidel and successive over relaxation are covered by the following result, where  $X \circ Y$  represents the Hadamard product of two matrices. Similarity is not restricted to the splittings covered by this theorem, for example, we can show that the result also holds for symmetric successive over-relaxation (SSOR).

**THEOREM 4.3** Suppose that  $A, H \in \mathbb{R}^{n \times n}$  and  $D$  is a nonsingular diagonal matrix. Let  $M = A \circ H$  and  $N = M - A$ . Then the spectrum of  $M^{-1}N$  is unchanged if we replace  $A$  with  $DAD$ .

*Proof.* Note that  $N = A \circ G$  where  $G = ee^T - H$  and so

$$(DAD \circ H)^{-1}(DAD \circ G) = (D(A \circ H)D)^{-1}D(A \circ G)D = D^{-1}(A \circ H)^{-1}(A \circ G)D. \quad \square$$

In particular, in the case of Gauss–Seidel, we get a bound comparable with the formulation adopted in [Livne & Golub \(2004\)](#) mentioned at the start of this section.  $R(x_*)$  is usually a little smaller for Livne and Golub’s method but for sparse matrices the difference is frequently negligible.

For SK though, we get  $R(x_*) = \rho(P) = 1$ . This is because the formulation described in (2.4) oscillates rather than converges. This phenomenon is discussed in [Knight \(2008\)](#), where it is shown that  $R(x_*)$  is equal to the modulus of the subdominant eigenvalue of  $P$ .

If  $A$  is symmetric, then it is possible to significantly improve convergence speed over SK with an appropriate choice of  $M$ . However, things are different if  $A$  is nonsymmetric. For example, consider the effect of using Gauss–Seidel in (4.1) where  $A$  is replaced by

$$S = \begin{bmatrix} 0 & A \\ A^T & 0 \end{bmatrix}$$

and  $x_k = [r_k^T \quad c_k^T]^T$ . We have

$$\begin{bmatrix} r_{k+1} \\ c_{k+1} \end{bmatrix} = \begin{bmatrix} r_k \\ c_k \end{bmatrix} + \begin{bmatrix} \mathcal{D}(Ac_k)\mathcal{D}(r_k)^{-1} & 0 \\ A^T & \mathcal{D}(A^T r_k)\mathcal{D}(c_k)^{-1} \end{bmatrix}^{-1} \begin{bmatrix} \mathcal{D}(r_k)^{-1}e - Ac_k \\ \mathcal{D}(c_k)^{-1}e - A^T r_k \end{bmatrix}.$$

After some straightforward manipulation, this becomes

$$\begin{bmatrix} r_{k+1} \\ c_{k+1} \end{bmatrix} = \begin{bmatrix} \mathcal{D}(Ac_k)^{-1}e \\ \mathcal{D}(A^T r_k)^{-1}e + c_k - \mathcal{D}(c_k)\mathcal{D}(A^T r_k)^{-1}A^T r_{k+1} \end{bmatrix}. \quad (4.3)$$

In the spirit of Gauss–Seidel, we can replace  $r_k$  on the right-hand side of (4.3) with  $r_{k+1}$ , giving

$$\begin{bmatrix} r_{k+1} \\ c_{k+1} \end{bmatrix} = \begin{bmatrix} \mathcal{D}(Ac_k)^{-1}e \\ \mathcal{D}(A^T r_{k+1})^{-1}e \end{bmatrix},$$

and this is precisely (2.2), that is, SK. In other words the Gauss–Seidel Newton method offers no improvement over SK for nonsymmetric matrices.

There are any number of choices for  $M$  in (4.1), but in tests we have not been able to gain a consistent and significant improvement over SK when  $A$  is nonsymmetric. As most of the applications for balancing involve nonsymmetric matrices we need a different approach.

## 5. Conjugate gradient method

Recall that if  $A$  is symmetric and non-negative, the Newton step (2.5) can be written

$$\mathcal{A}_k x_{k+1} = Ax_k + \mathcal{D}(x_k)^{-1}e,$$

where  $\mathcal{A}_k = A + \mathcal{D}(x_k)^{-1}\mathcal{D}(Ax_k)$ . By Theorems 3.1 and 3.5, we know that  $\mathcal{A}_k$  is positive semidefinite and (2.5) is consistent whenever  $x_k > 0$ , and we can solve the Newton step with the conjugate gradient method (Hestenes & Stiefel, 1952). Essentially, all we need to do is to find an approximate solution to (2.5) and iterate, ensuring that we never let components of our iterates become negative. We now look in more detail at how we implement the method efficiently.

First, motivated by the proof of Theorem 3.1, we note that we can apply a diagonal scaling to (2.5) to give a symmetric diagonally dominant system. Premultiplying each side of the equation by  $\mathcal{D}(x_k)$  and writing  $y_{k+1} = \mathcal{D}(x_k)^{-1}x_{k+1}$  we get

$$(B_k + \mathcal{D}(B_k e))y_{k+1} = (B_k + I)e, \quad (5.1)$$

where  $B_k = \mathcal{D}(x_k)A\mathcal{D}(x_k)$ . We need not form  $B_k$  explicitly, as all our calculations can be performed with  $A$ . The natural choice as initial iterate for every inner iteration is  $e$  (we also use  $e$  as the default choice to start the algorithm), for which the initial residual is

$$r = (B_k + I)e - (B_k + \mathcal{D}(B_k e))e = e - B_k e = e - v_k,$$

where  $v_k = x_k \circ Ax_k$ , and  $\circ$  refers to the Hadamard (elemental) product. Inside the conjugate gradient iteration we need to perform a matrix–vector product involving  $B_k + \mathcal{D}(B_k e)$ . For a given vector  $p$ , we can perform this efficiently using the identity

$$(B_k + \mathcal{D}(B_k e))p = x_k \circ (A(x_k \circ p)) + v_k \circ p.$$

Since the rest of the conjugate gradient algorithm can be implemented in a standard way, this matrix–vector product is the dominant factor in the cost of the algorithm as a whole.

As a stopping criterion for the scaling algorithm we use the residual measure

$$\|e - \mathcal{D}(x_k)Ax_k\|_2 = \|e - v_k\|_2.$$

In contrast to our experience with stationary iterative methods, it pays to run the inner iteration for a number of steps, and so we need a stopping criterion for the inner iteration, too. As is standard with inexact Newton methods, we do not need an accurate estimate while we are a long way from the solution and we adopt the approach outlined in Kelley (2003, Chapter 3) due to Dembo *et al.* (1982). To solve  $F(x) = 0$ , the authors advocate stopping the inner iteration when a bound of the form

$$\|F(x_k) + F'(x_k)s_k\| \leq \eta_k \|F(x_k)\|$$

is satisfied and letting  $x_{k+1} = x_k + s_k$ . In our case, this bound can be expressed as

$$\|(B_k + \mathcal{D}(B_k e))y_{k+1}\| \leq \eta_k \|(B_{k-1} + \mathcal{D}(B_{k-1} e))y_k\|.$$

The details of how  $\eta_k$  is reduced are described in Kelley (2003). We use the parameters  $\eta_{\max}$  to limit the size of the forcing term  $\eta_k$  and  $\gamma$  to limit  $\eta_{k+1}/\eta_k$ . In our algorithm the default values are  $\eta_{\max} = 0.1$ ,  $\gamma = 0.9$ .

While we know that  $B_k + \mathcal{D}(B_k e)$  is SPD if  $x_k > 0$ , we cannot guarantee that (5.1) will have a positive solution. Furthermore, we do not know that our Newton method will converge if our initial guess is a long way from the solution. We therefore apply a box constraint inside the inner iteration. We introduce a parameter  $\delta$  which determines how much closer to the edge of the positive cone we are willing to let our current iterate move and a second parameter  $\Delta$  bounding how far we can move away. We do not want components of our scaling to grow too large as this generally forces another component to shrink towards 0. By rewriting (2.5) in the form (5.1) we ensure that all coordinate directions are treated equally. Before we move along a search direction in the conjugate gradient algorithm, we check whether this would move our iterate outside our box. If it does, we can either reject the step or only move part of the way along it and then restart the inner iteration. In our experience, it pays not to reject the step completely and instead we move along it until the minimum element of  $y_{k+1}$  equals  $\delta$  or the maximum equals  $\Delta$ . In general, the choices  $\delta = 0.1$  and  $\Delta = 3$  seem to work well.

The method is a simplified version of the algorithm in O'Leary (1980) for solving constrained linear systems. We update the Newton step as soon as we hit the constraint rather than continuing with a partitioned iteration matrix. As we approach the solution to the balancing problem we can expect that (5.1) will have a solution  $y_{k+1} > 0$ , so eventually we will stop applying the box constraint.

We report the results of our tests of the method in the next section. The method appears to converge independently of the spectrum of our target matrix meaning that the conjugate gradient method converges much more quickly than the other methods we have discussed. Our experience is that it is not only quick but it is robust, too.

While stationary iterative methods could not improve on SK for nonsymmetric systems, the picture changes for the conjugate gradient method. If  $A$  is nonsymmetric we can apply the symmetric algorithm to

$$S = \begin{bmatrix} 0 & A \\ A^T & 0 \end{bmatrix}.$$

The resulting method is no longer connected to SK. From the results in Section 3 we know that the Jacobian will be singular at each step, but that the linear systems we solve will be consistent. Although  $B_k$  is singular, the fact that it is semidefinite and the system is consistent means the conjugate gradient method (CG) will work, provided the initial residual is in the orthogonal complement of the kernel of  $S$ . We ensure this is the case by setting  $y_0 = e$ . Our tests show our method is robust in these cases, although the residual need not decrease monotonically, suggesting that the algorithm may be moving between local minima of the function  $\|\mathcal{D}(x)Sx - e\|^2$ .

While the use of the augmented matrix  $S$  may seem inefficient, it is seemingly inevitable that we will have to work with both  $A$  and  $A^T$  in the algorithm in order to balance rows and columns simultaneously. Rather than working with  $S$ , we can unravel the method and work directly with  $A$  and  $A^T$ . This is particularly useful in solving the large-scale balancing problems discussed in Knight (2008), where a rank one correction is applied implicitly to  $A$ .

As we mentioned earlier, the practical application of the algorithm in Khachiyan & Kalantari (1992) is not discussed by its authors. However, it too can be implemented using the conjugate gradient method.

For this formulation, (5.1) is replaced with

$$(B_k + I)y_{k+1} = 2e.$$

While this is also straightforward to implement, we can no longer guarantee that the systems will be semidefinite unless  $A$  is too. In experiments, we have found that the Khachiyan–Kalantari approach is significantly slower than the one we propose.

## 6. Results

We now compare the performance of the conjugate-gradient-based approach against a number of other algorithms. The algorithms considered are as follows: BNEWT is our implementation of an inexact Newton iteration with conjugate gradients; SK is the Sinkhorn–Knopp algorithm and EQ is the ‘equalize’ algorithm from Parlett & Landis (1982); GS is a Gauss–Seidel implementation of (4.1). We use this in preference to the method outlined in Livne & Golub (2004) since for large matrices it is much easier to implement in MATLAB. In terms of rate of convergence, the algorithms are similar.

We have tested the algorithms for both symmetric and nonsymmetric matrices. All our tests were carried out in MATLAB. We measure the cost in terms of the number of matrix–vector products taken to achieve the desired level of convergence. In our tests on nonsymmetric matrices, we have counted the number of products involving  $A$  or  $A^T$  (double the number for the symmetrised matrix  $S$ ). We ran our algorithms until an appropriate residual norm fell below a tolerance of  $10^{-6}$ . For BNEWT, and other algorithms for symmetric matrices, we measured  $\|\mathcal{D}(x_k)Ax_k - e\|_2$  (replacing  $A$  with  $S$  if  $A$  was nonsymmetric). For SK we measured  $\|\mathcal{D}(c_k)A^T r_k - e\|_2$ , where  $c_k$  and  $r_k$  are defined in (2.2). In all cases, our initial iterate is a vector of ones.

In BNEWT there are a number of parameters ( $\eta_{\max}$ ,  $\gamma$ ,  $\delta$  and  $\Delta$ ) connected to the convergence criterion of the inner step and the line search that we can tune to improve performance. Unless otherwise stated, we have used the default values for these parameters.

Our first batch of test matrices were used in Parlett & Landis (1982) to compare their balancing algorithms to SK. These are all  $10 \times 10$  upper Hessenberg matrices defined as follows.  $H = (h_{ij})$  where

$$h_{ij} = \begin{cases} 0 & \text{if } j < i - 1, \\ 1 & \text{otherwise.} \end{cases}$$

$H_2$  differs from  $H$  only in that  $h_{12}$  is replaced by 100 and  $H_3 = H + 99I$ . Our results are given in Table 1. In this experiment, the tolerance changed to  $10^{-5}$  as this was the choice in Parlett & Landis (1982). EQ uses mainly vector–vector operations, and we have converted this to an (approximately) equivalent number of matrix–vector products.


We see the consistent performance of BNEWT, outperforming the other choices. The results for GS confirm our analysis in Section 4, showing it is virtually identical to SK. The more fine-grained nature

TABLE 1 *Number of matrix–vector products for Parlett and Landis test matrices*

|       | BNEWT | SK   | EQ  | GS   |
|-------|-------|------|-----|------|
| $H$   | 76    | 110  | 141 | 114  |
| $H_2$ | 90    | 144  | 131 | 150  |
| $H_3$ | 94    | 2008 | 237 | 2012 |

TABLE 2 *Rate of convergence for  $H_n$* 

|       | $n = 10$ | 25              | 50                 | 100                |
|-------|----------|-----------------|--------------------|--------------------|
| BNEWT | 124      | 300             | 660                | 1792               |
| SK    | 3070     | 16258           | 61458              | 235478             |
| $r_n$ | 217      | $7 \times 10^6$ | $2 \times 10^{14}$ | $2 \times 10^{29}$ |

FIG. 1. Convergence graph of BNEWT for  $H_{50}$ .

of EQ means that a comparison with other algorithms in terms of operation counts is misleading. In our tests, it took a similar number of iterations as SK did to converge, and a significantly longer time.

We next tested BNEWT on the  $n \times n$  version of  $H_3$ . For large values of  $n$ , this becomes very challenging to balance since the ratio between the smallest and largest elements of the balancing factors grows extremely large (the matrix becomes very close in a relative sense to a matrix without total support). The rate of convergence is given in Table 2, along with the ratio  $r_n = (\max_i x_i) / (\min_i x_i)$ .

BNEWT still copes in extremely trying circumstances, although the convergence of the residual is far from monotonic (we illustrate the progress after each inexact Newton step for  $n = 50$  in Fig. 1).

The oscillatory behaviour is undesirable, and it can be ameliorated somewhat by varying the parameters in BNEWT. In terms of cost, the choice  $\eta_{\max} = 10^{-2}$ ,  $\delta = 0.25$  proved best in reducing the matrix–vector product count to 568.

In our next test we look at three matrices from the Harwell–Boeing collection (Duff *et al.*, 1992), namely, GRE185, GRE343 and GRE1107 (the number representing the dimension). The smallest of these was considered in Livne & Golub (2004) as a candidate for preconditioning, but convergence was slow. We illustrate the progress of BNEWT in Fig. 2. The cost in matrix–vector products (in ascending order of  $n$ ) was 206, 118 and 350. Again, the method proves robust and we avoid the oscillatory behaviour we saw for our previous, more extreme, example.




FIG. 2. Convergence of BNEWT for sparse nonsymmetric matrices.

We have also carried out more comprehensive tests on a selection of matrices from the University of Florida Sparse Matrix Collection (UFL).<sup>1</sup> First we attempted to scale 45 sparse symmetric matrices with dimensions between 5,000 and 50,000 from the matrix sets `Schenk_IBMNA` and `GHS_indef` which have been observed to be resistant to scaling algorithms (Ruiz & Ucar, 2011). Each of the algorithms BNEWT, GS and SK was run until our default tolerance was reached or 2,000 matrix–vector products had been computed. In all but one case, BNEWT converged within the limit while GS met the convergence criterion in 25 cases, and SK in only 3. In Fig. 3 we plot the number of matrix–vector products required by BNEWT (along the  $x$ -axis) against the number required by GS. Cases where the algorithms failed to converge in time lie at the limits of the axes and are marked with triangles. The three cases where SK converged are circled (in each case requiring significantly more work than the other algorithms). A pattern emerges where GS converges very quickly on ‘easy’ examples but BNEWT is far more robust and almost always converges at reasonable speed.

Similar tests were carried out on 60 representative nonsymmetric matrices from the UFL collection with dimensions between 50 and 50,000 to compare BNEWT and SK. The results are shown in Fig. 4. In this case, the algorithms were given up to 50,000 matrix–vector products to converge. We see the clear superiority of BNEWT in these examples, often at least 10 times as fast as SK, although on occasions convergence was not particularly rapid, particularly when  $(\max_i x_i)/(\min_i x_i)$  is large. SK failed to converge in over half the cases.

We note that many of the sparse matrices in the UFL collection do not have total support, and we omitted these from the experiment. We tested BNEWT on a number of sparse nonsymmetric matrices which have support but lack total support. In these cases, balancing is impossible but an approximate solution can be found. In many cases BNEWT is robust enough to find an approximate solution in reasonable time (nonzero elements which do not lie on a positive diagonal are almost obliterated by the scaling). Such approximate scalings can be used as preconditioners, but we have found that equally

<sup>1</sup> Web address <http://www.cise.ufl.edu/research/sparse/matrices> (last accessed 10th September 2012).

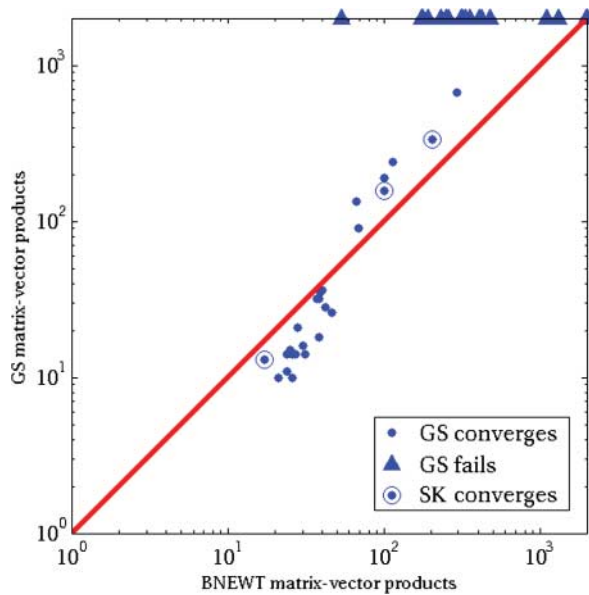


FIG. 3. Comparison of algorithms on sparse symmetric matrices.

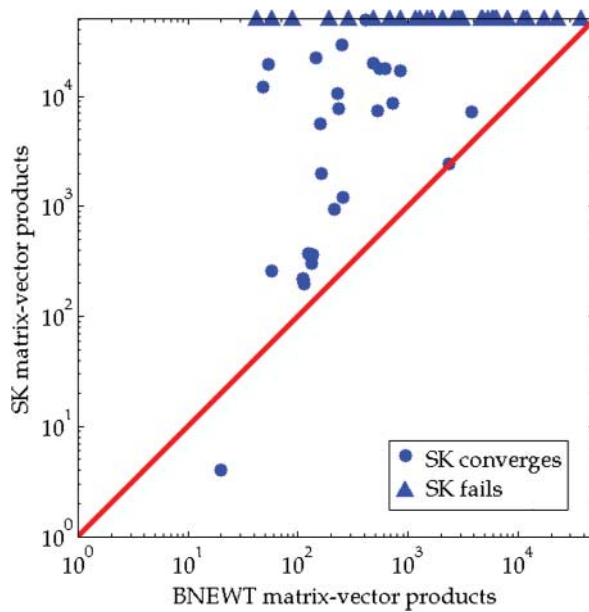


FIG. 4. Comparison of algorithms on sparse nonsymmetric matrices.



good preconditioners can be found by applying a handful of iterations of SK, even though the resulting diagonal factors do not come close to balancing the matrix.

## 7. Conclusions

In many applications, the Sinkhorn–Knopp algorithm offers a simple and reliable approach to balancing, but its limitations become clear when the matrices involved exhibit a degree of sparsity (in particular, the *a priori* measure of convergence given in Franklin & Lorenz (1989) fails if the matrix to be balanced has even a single zero). We have clearly shown that an inexact Newton method offers a simple way of overcoming these limitations. By taking care to remain in the positive cone, we see improvements for symmetric and nonsymmetric matrices. Further speed-ups may be possible by using more sophisticated preconditioning techniques for the conjugate gradient algorithm. In examples where our new algorithm converges slowly, the diagonal factors typically contain widely varying components ( $(\max_i x_i)/(\min_i x_i) \gg 10^{10}$ , say) and we are not aware of applications for accurate balancing in these cases.

## Funding

This work was supported in part by a grant from the Carnegie Trust for the Universities of Scotland.

## REFERENCES

- BACHARACH, M. (1970) *Biproportional Matrices & Input-Output Change*. Cambridge: Cambridge University Press.
- BALAKRISHNAN, A., HWANG, I. & TOMLIN, C. J. (2004) Polynomial approximation algorithms for belief matrix maintenance in identity management. *Proceedings of the 43rd IEEE Conference on Decision and Control*, **5**, (Editors) City: Publisher, pp. 4874–4879.
- BALINSKI, M. (2008) Fair majority voting (or how to eliminate gerrymandering). *Amer. Math. Monthly*, **115**, 97–113.
- BRADLEY, A. M. (2010) Algorithms for the equilibration of matrices and their application to limited-memory quasi-Newton methods. *Ph.D. Thesis*, Stanford University, USA.
- DASZYKOWSKI, M., MOSLETH FRGESTAD, E., GROVE, H., MARTENS, H. & WALCZAK, B. (2009) Matching 2D gel electrophoresis images with Matlab ‘Image Processing Toolbox’. *Chemometr. Intell. Lab. Syst.*, **96**, 188–195.
- DEMBO, R. S., EISENSTAT, S. C. & STEIHAUG, T. (1982) Inexact Newton methods. *SIAM J. Numer. Anal.*, **19**, 400–408.
- DUFF, I. S., GRIMES, R. G. & LEWIS, J. G. (1992) Users’ guide for the Harwell-Boeing sparse matrix collection (release I). *Technical Report RAL 92-086*. Rutherford Appleton Laboratory.
- FRANKLIN, J. & LORENZ, J. (1989) On the scaling of multidimensional matrices. *Linear Algebra Appl.*, **114/115**, 717–735.
- FÜRER, M. (2004) Quadratic convergence for scaling of matrices. *Proc. ALNEX and ANALC 2004* (L. Arge, G. F. Italiano & R. Sedgewick eds). City: Publisher pp. 216–223.
- HESTENES, M. R. & STIEFEL, E. (1952) Methods of conjugate gradients for solving linear systems. *J. Res. Natl. Bur. Stand.*, **49**, 409–436.
- KALANTARI, B. & KHACHIAN, L. (1996) On the complexity of nonnegative-matrix scaling. *Linear Algebra Appl.*, **240**, 87–103.
- KELLEY, C. T. (2003) *Solving Nonlinear Equations with Newton’s Method*. Philadelphia: SIAM.
- KHACHIAN, L. & KALANTARI, B. (1992) Diagonal matrix scaling and linear programming. *SIAM J. Opt.*, **2**, 668–672.
- KNIGHT, P. A. (2008) The Sinkhorn–Knopp algorithm: convergence and applications. *SIMAX*, **30**, 261–275.
- KOLOTILINA, L. YU. (2003) Nonsingularity/singularity criteria for nonstrictly block diagonally dominant matrices. *Linear Algebra Appl.*, **359**, 133–159.



- LAMOND, B. & STEWART, N. F. (1981) Bregman's balancing method. *Transp. Res. B*, **15**, 239–248.
- LINIAL, N., SAMORODNITSKY, A. & WIGDERSON, A. (2000) Deterministic strongly polynomial algorithm for matrix scaling and approximate permanents. *Combinatorica*, **20**, 545–568.
- LIVNE, O. E. & GOLUB, G. H. (2004) Scaling by binormalization. *Numer. Algorithms*, **35**, 97–120.
- MARSHALL, A. W. & OLKIN, I. (1968) Scaling of matrices to achieve specified row and column sums. *Numer. Math.*, **12**, 83–90.
- O'LEARY, D. P. (1980) A generalized conjugate gradient algorithm for solving a class of quadratic programming problems. *Linear Algebra Appl.*, **34**, 371–399.
- ORTEGA, J. M. & RHEINOLDT, W. C. (1970) *Iterative Solution of Nonlinear Equations in Several Variables*. New York: Academic Press.
- PARLETT, B. N. & LANDIS, T. L. (1982) Methods for scaling to doubly stochastic form. *Linear Algebra Appl.*, **48**, 53–79.
- RUIZ, D. & UCAR, B. (2011) A symmetry preserving algorithm for matrix scaling. *Technical report RR-7552*. INRIA.
- SCHNEIDER, M. H. (1990) Matrix scaling, entropy minimization, and conjugate duality (II): the dual problem. *Math. Prog.*, **48**, 103–124.
- SINKHORN, R. & KNOPP, P. (1967) Concerning nonnegative matrices and doubly stochastic matrices. *Pacific J. Math.*, **21**, 343–348.
- SOULES, G. W. (1991) The rate of convergence of Sinkhorn balancing. *Linear Algebra Appl.*, **150**, 3–40.
- TAUSSKY, O. (1948) Bounds for characteristic roots of matrices. *Duke Math. J.*, **15**, 1043–1044.

## Appendix. The symmetric algorithm

```
function [x,res] = bnewt(A,tol,x0,delta,Delta,fl)
% BNEWT A balancing algorithm for symmetric matrices
%
% X = BNEWT(A) attempts to find a vector X such that
% diag(X)*A*diag(X) is close to doubly stochastic. A must
% be symmetric and nonnegative.
%
% X0: initial guess. TOL: error tolerance.
% delta/Delta: how close/far balancing vectors can get
% to/from the edge of the positive cone.
% We use a relative measure on the size of elements.
% FL: intermediate convergence statistics on/off.
% RES: residual error, measured by norm(diag(x)*A*x - e).

% Initialise
n = size(A,1); e = ones(n,1); res=[];
if nargin < 6, fl = 0; end
if nargin < 5, Delta = 3; end
if nargin < 4, delta = 0.1; end
if nargin < 3, x0 = e; end
if nargin < 2, tol = 1e-6; end
% Inner stopping criterion parameters.
g=0.9; etamax = 0.1;
eta = etamax; stop_tol = tol*.5;

x = x0; rt = tol^2; v = x.*(A*x); rk = 1 - v;
```

```

rho_kml = rk'*rk; rout = rho_kml; rold = rout;

MVP = 0; % We'll count matrix vector products.
i = 0; % Outer iteration count.

if fl == 1, fprintf('it in. it res\n'), end
while rout > rt % Outer iteration
    i = i + 1; k = 0; y = e;
    innertol = max([eta^2*rout,rt]);

    while rho_kml > innertol %Inner iteration by CG
        k = k + 1;

        if k == 1
            Z = rk./v; p=Z; rho_kml = rk'*Z;
        else
            beta=rho_kml/rho_km2;
            p=Z + beta*p;
        end

        % Update search direction efficiently.
        w = x.*(A*(x.*p)) + v.*p;
        alpha = rho_kml/(p'*w);
        ap = alpha*p;

        % Test distance to boundary of cone.
        ynew = y + ap;
        if min(ynew) <= delta
            if delta == 0, break, end
            ind = find(ap < 0);
            gamma = min((delta - y(ind))./ap(ind));
            y = y + gamma*ap;
            break
        end
        if max(ynew) >= Delta
            ind = find(ynew > Delta);
            gamma = min((Delta-y(ind))./ap(ind));
            y = y + gamma*ap;
            break
        end
        y = ynew;
        rk = rk - alpha*w; rho_km2 = rho_kml;
        Z = rk./v; rho_kml = rk'*Z;
    end

    x = x.*y; v = x.*(A*x);
    rk = 1 - v; rho_kml = rk'*rk; rout = rho_kml;
    MVP = MVP + k + 1;

```

```
% Update inner iteration stopping criterion.
rat = rout/rold; rold = rout; res_norm = sqrt(rout);
eta_o = eta; eta = g*rat;
if g*eta_o^2 > 0.1
eta = max([eta,g*eta_o^2]);
end
eta = max([min([eta,etamax]),stop_tol/res_norm]);

if fl == 1
fprintf('%3d %6d %.3e \n', i,k, r_norm);
res=[res; r_norm];
end
end
fprintf('Matrix-vector products = %6d\n', MVP)
```